




Today

- Informed Search
 - Heuristics
 - Greedy Search
 - A* Search
- Graph Search



Recap: Search



Recap: Search

- Search problem:
 - States (configurations of the world)
 - Actions and costs
 - Successor function (world dynamics)
 - Start state and goal test


- Search tree:
 - Nodes: represent plans for reaching states
 - Plans have costs (sum of action costs)

- Search algorithm:
 - Systematically builds a search tree
 - Chooses an ordering of the frontier (unexplored nodes)
 - Optimal: finds least-cost plans



Uniform Cost Search

- Strategy: expand lowest path cost



- The good: UCS is complete and optimal!

- The bad:
 - Explores options in every "direction"
 - No information about goal location



Video of Demo Contours UCS Empty



Video of Demo Contours UCS Pacman Small Maze



Informed Search



What we would like to have happen

Guide search *towards the goal* instead of *all over the place*



Informed



Uninformed

Heuristic:


-

- a function that estimates how close the current state is to a goal.
- designed for a particular search problem

$n(i)$ = straight-line distance to Buchanan


$h(n)$ = straight-line distance to Bucharest

- $h(n)$ = Manhattan distance = $|\Delta x| + |\Delta y|$
- Is Manhattan better than straight-line distance?



Greedy Search

- Expand the node that seems closest...



- What can go wrong?

Everything is based on how good the heuristic is


Greedy Search

- Strategy: expand a node that you think is closest to a goal state
- Heuristic: estimate of distance to nearest goal for each state

- A common case:

- Best-first takes you **straight to the (wrong) goal**

- Worst-case: like a badly-guided DFS



Sometimes, an idiot, other times a servant

- Essentially, DFS (or BFS) with PQ

Video of Demo Contours Greedy (Empty)



Video of Demo Contours Greedy (Pacman Small Maze)




A* Search



* → optimal

Combining UCS and Greedy

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$




- A* Search orders by the sum: $f(n) = g(n) + h(n)$

backward cost + forward cost

Example: Teg Grenager

Is A* Optimal?



What went wrong?

- Actual bad solution cost < estimated good solution cost
- We need estimates to be less than actual costs!

Optimal:


Lowest cost among all solutions

Our heuristic SUCKS!


our estimates need to be less than or equal to actual costs

When should A* terminate?

- Should we stop when we enqueue a goal?



- No: only stop when we dequeue a goal



→ expand B


enqueue, pull next thing off

Stop when you dequeue the goal, not enqueue it

Admissible Heuristics



Idea: Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the frontier.

Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs.

Snitty

Less Shitty

A "bad" heuristic will break optimality by trapping good plans on the frontier.

UNDERESTIMATE

Admissible Heuristics

- A heuristic h is **admissible** (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

- Examples:



- Coming up with admissible heuristics is most of what's involved in using A* in practice.
 - As we will talk about later, these should also be fast to compute.

if my heuristic guess is positive and less than the true cost to the nearest goal. In other words, a heuristic is optimistic if it's positive and less than the actual true cost. NOT OKAY w/ OVERESTIMATING

15 is closer to true cost

Optimality of A* Tree Search



Optimality of A* Tree Search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

- A will be chosen for expansion before B
- In other words, A will exit the frontier before B



Frontier



Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (maybe A itself)

Claim: n will be expanded before B

- $f(n)$ is less than or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f -cost
Admissibility of h
 $h = 0$ at a goal


Optimality of A* Tree Search: Blocking

Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (maybe A itself)

Claim: n will be expanded before B

- $f(n)$ is less than or equal to $f(A)$
- $f(A)$ is less than $f(B)$



$$g(A) < g(B)$$

$$f(A) < f(B)$$

Suboptimality of B
 $h = 0$ at a goal

Optimality of A* Tree Search: Blocking


Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (maybe A itself)

Claim: n will be expanded before B

- $f(n)$ is less than or equal to $f(A)$
- $f(A)$ is less than $f(B)$
- n is expanded before B

- All ancestors of A are expanded before B
- A is expanded before B
- A* tree search is optimal



$$f(n) \leq f(A) < f(B)$$

backward

$$f(n) = g(n) + h(n)$$

$$① f(n) \leq f(A)$$

$$② f(A) \leq f(B)$$

$$③ \therefore n \text{ exits before } B$$

Proof:

$$f(n) = g(n) + h(n)$$

$$f(A) = g(A) + h(A)$$

$$f(A) = g(A) \leftarrow \text{b.c. } A \text{ is an optimal goal}$$

$$g(n) \leq g(A) \Rightarrow h(n) \leq g(A) - g(n)$$

actual cost to n

actual cost to n

\rightarrow this is the true cost of n to A . $h(n)$ has to UNDER estimate this distance!

$$f(n) \leq g(n) + g(A) - g(n)$$

$$f(n) \leq g(A)$$

backward cost

forward

$$f(n) = g(n) + h(n)$$

$$f(A) = g(A) + h(A)$$

$$f(A) = g(A) + 0 \leftarrow$$

$$\Rightarrow f(A) = g(A)$$

$$f(n) = g(n) + h(n)$$


$$\leq g(A)$$

$$\leq f(A)$$

$$g(A) < g(B)$$

$$f(A) < f(B)$$

b/c h is admissible



\leq the true cost

UCS vs A* Contours

- Uniform-cost expands equally in all "directions"



- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Prove A* optimal

$$f(n) = g(n) + h(n)$$

$$f(A) = g(A) + h(A)$$

$$f(A) \leq g(A) + 0$$

$$\Rightarrow f(A) = g(A)$$

$$g(n) \leq g(A) \Rightarrow h(n) \leq g(A) - g(n)$$

↳ under estimate!



$$f(n) = g(n) + h(n)$$

$$f(A) = g(A) + h(A)$$

$$f(A) = g(A)$$

$$f(A) \leq g(A)$$

$$f(A) = g(A)$$

$$g(n) \leq g(A)$$

$$\therefore f(n) \leq f(A) < f(B)$$

Comparison



Greedy (h)




Uniform Cost (g)



A* (g+h)

A* Applications

- Video games
- Pathing / routing problems
- Resource planning problems
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition
- Protein design
- Chemical synthesis
- ...




Creating Heuristics



Creating Admissible Heuristics


- Most of the work in solving hard search problems optimally is in finding admissible heuristics
- Often, admissible heuristics are solutions to **relaxed problems**, where new actions are available



- Problem P_2 is a relaxed version of P_1 if $\mathcal{A}_2(s) \supseteq \mathcal{A}_1(s)$ for every s
- Theorem: $h_2^*(s) \leq h_1^*(s)$ for every s , so $h_2^*(s)$ is admissible for P_1
- Inadmissible heuristics are often useful too (will discuss later why!)

Take some of the "restriction" and throw them away
 → example: removing walls in manhattan
 → relaxed problem = under estimate

Example: 8 Puzzle



Start State



Actions



Goal State

- What are the states?
- How many states?
- What are the actions?
- What are the step costs?

Heuristic → how far am I from the thing

→ relaxed problem: manipulate however you want and make heuristic

8 Puzzle I

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = 8$



Start State

Goal State



Average nodes expanded when the optimal path has...			
...4 steps	...8 steps	...12 steps	
UCS	112	6,300	3.6×10^6
A* TILES	13	39	227

Statistics from Andrew Moore

8 Puzzle II

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?



Start State

Goal State

- Total Manhattan distance

- Why is it admissible?

- $h(\text{start}) = 3 + 1 + 2 + \dots = 18$

Average nodes expanded when the optimal path has...			
...4 steps	...8 steps	...12 steps	
A* TILES	13	39	227
A* MANHATTAN	12	25	73

8 Puzzle III

- How about using the *actual cost* as a heuristic?

- Would it be admissible?
- Would we save on nodes expanded?
- What's wrong with it?



- With A*: a trade-off between quality of estimate and work per node
- As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

Semi-Lattice of Heuristics


Trivial Heuristics, Dominance

- Dominance: $h_a \geq h_b$ if $\forall n : h_a(n) \geq h_b(n)$
 - Basically, this is a fancy way of saying that if you are bigger, you can assert dominance. Just like in real life.
- What if no one can assert dominance?
 - Create a super strain of dominance... is it admissible?
- Heuristics form a semi-lattice: $h(n) = \max(h_a(n), h_b(n))$
 - See the picture!
- Trivial heuristics
 - Bottom of lattice is the zero heuristic (what does this give us?)
 - Top of lattice is the exact heuristic




Find the silver lining

Graph Search




Tree Search: Extra Work!

- Failure to detect repeated states can cause exponentially more work.



Graph Search

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)



Graph Search

- Idea: never expand a state twice

- How to implement:

- Tree search + set of expanded states ("closed set")
- Expand the search tree node-by-node, but...
- Before expanding a node, check to make sure its state has never been expanded before
- If not new, skip it, if new add to closed set

- Important: store the closed set as a set, not a list

- Can graph search wreck completeness? Why/why not?

- How about optimality?


issue = we have to store every single state

Yes


Yes,

A* Graph Search Gone Wrong?


State space graph



Search tree



Consistency of Heuristics



- Main idea: estimated heuristic costs \leq actual costs
 - Admissibility: heuristic cost \leq actual cost to goal
 - $h(A) \leq \text{actual cost from A to G}$
 - Consistency: heuristic "arc" cost \leq actual cost for each arc
 - $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
- Consequences of consistency:
 - The f value along a path never decreases
 - $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
 - A* graph search is optimal


Optimality of A* Graph Search



Optimality of A* Graph Search


- Sketch: consider what A* does with a consistent heuristic:

- Fact 1: In tree search, A* expands nodes in increasing total f value (f-contours)
- Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally
- Result: A* graph search is optimal



Optimality

- Tree search:
 - A* is optimal if heuristic is admissible
 - UCS is a special case ($h = 0$)
- Graph search:
 - A* optimal if heuristic is consistent
 - UCS optimal ($h = 0$ is consistent)
- Consistency implies admissibility
- Most natural admissible heuristics tend to be consistent, especially if from relaxed problems



A*: Summary



A*: Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems

