

• Moore's Law: computing speed doubles every 18 months

↳ For this to continue, we must use Parallel computing
Essentially, it's about dividing the work among multiple CPUs

↳ No good programming languages for parallel.

Types of Parallel computing:

- Shared memory multiprocessors (Chips)
- ↳ OpenMP/threads

- GPU (NVIDIA)
- ↳ CUDA: coding environment for GPUs

- Distributed Parallel (cluster)
- ↳ Big Data, MPI

Big-O Analysis

$n = \text{limit size}$

- Analyze Small primes Sieve

- 1^{st} for loop $\rightarrow O(n)$

- 2^{nd} for loop $\rightarrow O(n)$

- 3^{rd} for loop $\rightarrow O(n)$

- Big-OH \rightarrow ignore lower order terms,

for ($d=2, d \leq n, d++$) $\leftarrow O(\sqrt{n})$

if ($i+d \geq 3$)

[for ($q=2d, q \leq n, q=q+d$) $\leftarrow \frac{n}{d}$
 $i+d = \text{false}$]

$$= \sum_{\substack{\text{Primes} \\ d \leq n}} \frac{n}{d} = n \left(\sum_{\substack{\text{Primes} \\ d \leq n}} \frac{1}{d} \right) O(\log \log n)$$

$= O(n \log \log n)$

To find primes $\leq n$, it takes time
 $O(n \log \log n)$

If we use P processors how long
Should it take?

$\hookrightarrow O\left(\frac{n \log \log n}{P}\right)$

How long does it take MPI to
get our processors ready?
 $P = \text{processors}$

$\hookrightarrow \text{MPI_INIT} \rightarrow O(\log P)$

$\rightarrow \text{MPI_Broadcast} \rightarrow O(\log P)$

$\rightarrow \text{MPI_Allreduce} \rightarrow O(2 \cdot \log P)$
or $O(1.5 \cdot \log P)$

$\rightarrow \text{MPI_Finalize} \rightarrow O(\log P)$

function findSmallPrimes(Sartn)

$\hookrightarrow O(\sqrt{n} \log \log n)$

Blocks are size \sqrt{n}

There are $\frac{n}{\sqrt{n}} = \sqrt{n}$ blocks

Each processor does $\frac{\sqrt{n}}{P}$ blocks

Riemann Zeta Function

$$\zeta(s) = \frac{1}{1^s} + \frac{1}{2^s} + \frac{1}{3^s} + \dots = \sum_{n=1}^{\infty} \frac{1}{n^s}$$



$$\text{csize} = \lceil \frac{\text{limit}}{P} \rceil$$

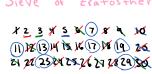
$$\text{Start} = 1 + (\text{csize} - 1) \cdot \text{csize}$$

$$\text{Stop} = \text{Start} + \text{csize} - 1$$

3. Share primes forward

Finding Primes

Sieve of Eratosthenes



Target n

1. Find primes $\leq \sqrt{n}$

2. Setipe Segments of size near \sqrt{n}

A processor makes an array or buffer of

cross off multiples of the primes $\leq \sqrt{n}$

3. Share primes forward

Maximum

given a list $A[0] \dots A[n-1]$
Find the max value.

Like the sum problem,
 $O(\log n)$ time, $O(n/\log n)$ processors

EREW PRAM

$O(1)$ Time!! But how?

initialize $B[i] = 1$

Proc. $i, j \neq i$

if $A[j] > A[i]$ then $B[i] = 0$

if $B[i] = 1$

max = $B[i]$

NO LOOP

$O(n^2)$ work but $O(1)$ run time

CRCW PRAM

Common

PRAM Model:



The processors work in
"Lockstep"

n processors for EACH $A[i]$

↳ initialize $B[i] = 1$

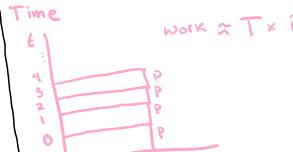
Proc. $i, j \neq i$

if $A[j] > A[i]$ then $B[i] = 0$

if $B[i] = 1$

max = $B[i]$

Time



work $\approx T \times P$

Goal of Designing Parallel Program:

Work $\approx T \times P$

Parallel Random Access Machine:

PRAM Flavors:

- EREW \rightarrow exclusive read exclusive write

- CREW \rightarrow Concurrent read exclusive write

- CRCW \rightarrow concurrent read and write

↳ Common: must write same value

↳ Arbitrary: can't predict which one wins

↳ Priority: small ID # wins

Hardware Power increases

PRAM:

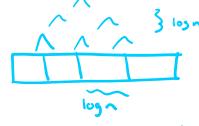
→ Adding up a list of numbers

$O(\log n)$ time

$O(n/\log n)$ processors

EREW PRAM

Work = time \times proc = $O(n)$
goal is for work to match sequential



If we have $P \leq O(n/\log n)$
then time = $\frac{\text{work}}{P} + \text{parallel time}$

$\hookrightarrow O(n/P + \log n)$

Maximum

- Crazy $O(1)$ time $O(n^2)$ processors

CRCW PRAM

for (int i = 0; i < n; i++)

 B[i] = 1

 for (int j = i+1; j < n; j++)

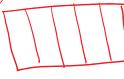
 if A[i] > A[j] then B[i] = 0

 if B[i] = 1 then output A[i]

Leveraging



apply crazy alg to there $\frac{n}{P}$ items $O(n)$ proc, $O(1)$ time



$O(\sqrt{n})$ processors, $O(1)$ times

↳ total $O(n)$ proc, $O(1)$ time

n pieces size \sqrt{n}
 $n^{3/4}$ pieces size $n^{1/4}$
 \vdots
 n pieces size 1

Total: $O(n)$ processors $O(1/\log n)$ time

One More Trick:

→ give each P chunks of size $n/\log n$ to final maximum sequentially

$\hookrightarrow O(\log n)$ time $O(n/\log n)$ processors

→ then, do the AsySym algorithm on

input size $n/\log n$.

Same time \leq proc.

Parallel Matrix Multiplication:

Recall:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} -1+0+3 \\ -4+0+6 \\ -7+0+9 \end{bmatrix} = \begin{bmatrix} 2 \\ 2 \\ 2 \end{bmatrix}$$

How do we parallelize for $n \times n$

→ Stride the rows

→ PRAM $\rightarrow n^2$ processors $\rightarrow O(n \log n)$ time

PRAM

Complete a sum of an array of length n

$$\begin{array}{ccccccccc} A & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \\ & [1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9] \end{array}$$



6 =

Matrix - Matrix Parallel

$$[n \times n] [n \times n] = [n \times n]$$

Normal Sequential Time: $O(n^3) \rightarrow n^3 \cdot n$

On the PRAM:

n^2 copies of dot prod

$O(\log n)$ time or $O(\log n)^2$ time

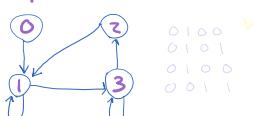
$O(n/\log n)$ space CREW PRAM

Graph Theory:

adjacency matrix

0	1	2	3
0	0	0	1
0	0	0	1
2	0	0	0
3	0	1	0

Reachability Graph distance = 2



For more example

$$A^k = R^{(k)}$$

$$B[i,j] = \sum_k A[i,k] \cdot A[k,j]$$

entry (i,j) of A^k

= number of times a

path of length k from i to j

Let R^k be the reachability graph for pairs of length $\leq k$

for a matrix M , let $b(M)$ be the matrix M with all non-zero entries set to 1

$$Then R^k = b(A + A^2 + A^3 + \dots + A^k)$$

A graph G is connected if its reachability graph R

has 1's everywhere (except possibly diagonal)

Computing A^k

$$A^{12} = 8 + 4$$

Matrix exponentiation \rightarrow binary algorithm to compute adjacency matrices

Fox's algorithm \rightarrow multiplying matrices that are large

Matrix Exponentiation:

idea:

to compute A^n

we compute $A, A^2, A^4, \dots, A^{2^{\lfloor \log n \rfloor}}$

then we use the pieces we

need based on writing n in

binary

Example:

$$A^8 = A^8 \cdot A^2 \cdot A^1$$

if in binary is $8 + 2 + 1$

1011

$$Y = I \quad n = 1011 \quad B = A$$

$$Y = Y \cdot B \quad n = n/2 \quad B = B^2 (= A^2) \quad \begin{matrix} 101 \\ = 101 \\ = \end{matrix}$$

$$Y = Y \cdot B \quad n = n/2 \quad B = B^2 (= A^4)$$

$$Y \text{ unchanged} \quad n = n/2 \quad B = B^2 (= A^8)$$

$$Y = Y \cdot B \quad n = n/2 \quad B = B^2 (= A^{16})$$

...

Each processor does $\frac{\sqrt{n}}{P}$ blocks

$$\sum_{\substack{\text{Primes} \\ d \leq \sqrt{n}}} \frac{\sqrt{n}}{d} = O(\sqrt{n} \log \log \sqrt{n})$$

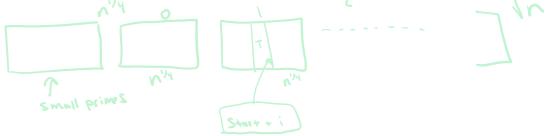
Time for 1 block

$$\rightarrow O\left(\frac{n \log \log n}{P} + \log P + \sqrt{n} \log \log n\right)$$

implies that $P \leq \sqrt{n}$

lower order terms

n factoring



$$n = 10000$$

$$n = 100$$

$$\sqrt{100} = 10$$

$$it = [1 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$$

for (int d=2;



$O(\log n)$ running time
 $O(n)$ processors

$O(\log n)$ time
 $O(n/\log n)$ processors
 $O(n)$ work
Aim: Work should match the Sequential run time
How do we do this?

1. each of the processors Sequential

\rightarrow PRAM \rightarrow in parallel \rightarrow $O(\log n)$ time

PRAM

Compute a sum of an array of length n
↳ done in $O(\log n)$ time (tree structure)
↳ $O(n \log n)$ processors
Work = $(\text{time} \cdot P) \rightarrow O(n)$

Dot Product

arrays a, b of length n ,

$$\text{Compute } \sum_{i=0}^{n-1} a[i] \cdot b[i]$$

Same: $O(\log n)$ time
 $O(n \log n)$ processor
EREW

Matrix - Vector product

$$[n \times n] [n] = [n]$$

is possible!
↳ $O(\log n)$ time
↳ $O(n \log n)$ processors

Simple n -way copy or dot product receives concurrent reads CREW

Before alg starts, Copy the vector n times:
Loop each vector entry into binary tree and copy down root to leaf in $O(\log n)$ time, n^2 processors

Then do logn factor trick in summation problem



Exponen

$y = I$

$B = A$

while

$i < 2$

3

$B = B \cdot A$

3

$A = A^2$

For

L

A

For

rewritten

$$Y = Y \cdot B \quad n = n/2 \quad B = B^2 \text{ (} : A^{16} \text{)}$$

$$Y = A^n$$

ntiation :

Example :

$$A^{20} \quad n = 16 + 4 = 10100$$

$$Y = I$$

$$B = A \quad n = 20, \text{ even}$$

$$B = B^2 \quad n = n/2 = 10$$

$$B = B^2 \cdot (A^n) \quad n = n/2 = 5$$

$$Y = Y \cdot B \quad (\text{so } Y = A^4)$$

$$B = B^2 \cdot (A^4) \quad n = n/2 = 2 \quad \text{even}$$

$$B = B^2 \cdot (A^4) \quad n = n/2 = 1 \quad \text{odd}$$

$$Y = Y \cdot B \quad (= A^4 \cdot A^4)$$

is in Y

z's Algorithm

When matrices are too big, we gotta use Fox!

- $B = C$ assume they are $n \times n$

$$\begin{matrix} A \\ \boxed{i} \\ j \end{matrix} \quad \begin{matrix} B \\ \boxed{i} \\ j \end{matrix} = \begin{matrix} C \\ \boxed{i} \\ j \end{matrix}$$

```
- (int i=0; i < n; i++)
for (int j=0; j < n; j++)
```

$$C[i, j] = 0;$$

```
For (int k=0; k < n; k++)
    C[i, j] += A[i, k] * B[k, j]
```

Ex:

$$\begin{matrix} A \\ \left[\begin{matrix} 1 & 0 & 1 \\ 1 & 1 & 0 \\ 0 & 1 & 1 \end{matrix} \right] \end{matrix} \quad \begin{matrix} B \\ \left[\begin{matrix} 1 & 2 & 3 \\ 3 & 2 & 1 \\ 0 & 1 & 0 \end{matrix} \right] \end{matrix} = \begin{matrix} C \\ \left[\begin{matrix} ? & ? & ? \\ ? & ? & ? \\ ? & ? & ? \end{matrix} \right] \end{matrix}$$

$P_{i,j}$ is the processor computing $C[i, j]$

→ Old version (Fox free)

We would have:

$$P_{00} A_{00} \cdot B_{00} \quad P_{01} A_{00} \cdot B_{01} \quad P_{02}$$

$$P_{10} A_{10} \cdot B_{00} \quad P_{11} \quad P_{12}$$

$$P_{20} A_{20} \cdot B_{00} \quad P_{21} \quad P_{22}$$

Fox Version:

$$P_{00} A_{00} \cdot B_{00} \quad P_{01} \quad P_{02}$$

$$P_{10} A_{10} \cdot B_{10} \quad P_{11} \quad P_{12}$$

$$P_{20} A_{20} \cdot B_{20} \quad P_{21} \quad P_{22}$$